

Intro to R, Markdown and the Tidyverse

Jameson Watts, Ph.D.

Agenda

1. The R programming environment
2. RMarkdown
3. Intro to Tidyverse

Strong suggestion: download some [cheat sheets](#)

The R programming environment

What is R?

- R is a software environment for statistical computing and graphics.
- It is also a programming language
 - Variables (of different types)
 - Structures
 - Functions
 - Syntax

Using RStudio

- *Console*: where you can execute code
- *Source file*: create and save scripts, which get sent to the console
- *Comments*: document what you are doing
- *Environment*: a snapshot of the current state of affairs
- *Help*: lookup how to do stuff
- *Packages*: bundles of code written by others

Exercise

Open RStudio and:

1. Create a new R script
2. Create a comment that explains what you are doing
3. Compute the solution to $2 \times (28 - 7)$
4. Run the script

Solution

```
# Math is fun!  
2*(28-7)
```

```
## [1] 42
```

Exercise

1. Edit your R script to assign the solution to $2 \times (28 - 7)$ to a variable
2. Multiply your variable by 42 and then take the square root
3. Display the answer

Hint: use `sqrt ()`

Solution

```
# Math is fun!  
s <- 2*(28-7)  
s <- s*42  
sqrt(s)
```

```
## [1] 42
```

Functions

Input → Function → Output

```
# Exponentiate a number and assign to a variable  
s <- exp(42)  
s
```

```
## [1] 1.739275e+18
```

```
# Take the log of this variable and display it  
log(s)
```

```
## [1] 42
```

Packages

- A package is a collection of functions, documentation, and sometimes data
- There are a number of packages that are part of base R
- You can install other packages from CRAN
- Not all packages are created equal

Note: The functions **exp()** and **log()** are part of Base R

Excercise

1. Install the tidyverse package
2. load tidyverse into memory

Solution

```
# install.packages('tidyverse')  
library(tidyverse)
```

```
## — Attaching packages
```

```
## ✓ ggplot2 3.2.1    ✓ purrr  0.3.2  
## ✓ tibble  2.1.3    ✓ dplyr  0.8.3  
## ✓ tidyr   1.0.0    ✓ stringr 1.4.0  
## ✓ readr   1.3.1    ✓ forcats 0.4.0
```

```
## — Conflicts
```

```
## ✖ dplyr::filter() masks stats::filter()  
## ✖ dplyr::lag()     masks stats::lag()
```

tidy

RMarkdown

The what and why of RMarkdown

RMarkdown is a file format that allows you to embed your code directly within your report.

Why is this important?

Preamble

To begin, you need to specify the type of document you want and set some options...

...you need to write some YAML

```
---  
title: "Intro to R, Markdown and the Tidyverse"  
author: "Jameson Watts, Ph.D."  
output:  
  ioslides_presentation:  
    widescreen: yes  
---
```

```
{r setup, include=FALSE}  
knitr::opts_chunk$set(echo = TRUE)
```

...but luckily this is all pretty much done for you when you create a new RMarkdown file in RStudio.

Formatting

Plain text

```
# header 1
## header 2
### header 3 etc.
```

an equation $y = mx + b$

image: ![caption]("pathtofile.png")

url: [link text](http://address.com)

> block quote

italics

****bold****

Formatting cont'd

superscript²

- unordered list
- next item
 - + some subitems
 - + subitem 2

1. unordered list
2. next item
 - + some subitems
 - + subitem 2

Table header	Column 2 header
cell 1	cell 2
cell 3	cell 4

Code Evaluation

Chunk it

```
` `` {r eval=FALSE, cache=FALSE}
```

```
# here is a comment
```

```
my_var <- 42
```

```
my_var
```

```
` ``
```

Inline

My favorite variable is: ``r my_var``

Code Chunk Options

- eval
- echo
- warning
- error
- message
- cache
- fig.width
- fig.height

...and so much more can be found [here](#)

Exercise: Reproduce the output of this slide

Some rad markdown

Note: Using **markdown** is an **easy** way to marry analysis with report writing.

The advantages are:

1. Simple, output-agnostic formatting
2. Reproducible results
3. Easy for others to follow your analysis

Example analysis:

```
my_var <- exp(log(42))
```

*Note*²: It is also cool because I can embed R code inline. For instance, the value of my_var is 42.

Intro to the Tidyverse

Philosophy

An 'opinionated' collection of R packages designed for data science. All packages share an underlying design philosophy, grammar, and data structures.

- readr – importing data
- dplyr – manipulating data
- tidyr – cleaning data
- ggplot2 – visualizing data

...more are being added regularly

Importing data (pair up and follow along)

1. Goto jamesonwatts.github.io, scroll down to the teaching section and click on the GSMDS Class link
2. Find the resources folder and download the [wine-data.csv](#) file
3. Create a new R code chunk in your RMD file and add the following:

```
library(tidyverse)
wine <- read_csv("../resources/winemag-data.csv")
```


Data Frames

```
wine <- read_csv("../resources/winemag-data.csv")
```

```
## Parsed with column specification:
## cols(
##   x1 = col_double(),
##   country = col_character(),
##   description = col_character(),
##   designation = col_character(),
##   points = col_double(),
##   price = col_double(),
##   province = col_character(),
##   region_1 = col_character(),
##   region_2 = col_character(),
##   taster_name = col_character(),
##   taster_twitter_handle = col_character(),
##   title = col_character(),
##   variety = col_character(),
##   winery = col_character()
## )
```

Rectangular data

- All columns are variables
- All rows are observations
- Each cell is a value

```
head(wine,5)
```

```
## # A tibble: 5 x 14
##       X1 country description designation points price province region_1
##   <dbl> <chr>   <chr>         <chr>         <dbl> <dbl> <chr>   <chr>
## 1     0 Italy   Aromas inc... Vulkà Bian...     87    NA Sicily ... Etna
## 2     1 Portug... This is ri... Avidagos       87    15 Douro   <NA>
## 3     2 US      Tart and s... <NA>           87    14 Oregon  Willame...
## 4     3 US      Pineapple ... Reserve La...     87    13 Michigan Lake Mi...
## 5     4 US      Much like ... Vintner's ...     87    65 Oregon  Willame...
## # ... with 6 more variables: region_2 <chr>, taster_name <chr>,
## #   taster_twitter_handle <chr>, title <chr>, variety <chr>, winery <chr>
```

Note: you can also open data files within RStudio just like in Excel

Let's take a glimpse instead

```
glimpse(wine)
```

```
## Observations: 129,971
## Variables: 14
## $ X1          <dbl> 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, ...
## $ country     <chr> "Italy", "Portugal", "US", "US", "US", "Sp...
## $ description <chr> "Aromas include tropical fruit, broom, bri...
## $ designation <chr> "Vulkà Bianco", "Avidagos", NA, "Reserve L...
## $ points      <dbl> 87, 87, 87, 87, 87, 87, 87, 87, 87, 87, 87...
## $ price       <dbl> NA, 15, 14, 13, 65, 15, 16, 24, 12, 27, 19...
## $ province    <chr> "Sicily & Sardinia", "Douro", "Oregon", "M...
## $ region_1    <chr> "Etna", NA, "Willamette Valley", "Lake Mic...
## $ region_2    <chr> NA, NA, "Willamette Valley", NA, "Willamet...
## $ taster_name <chr> "Kerin O'Keefe", "Roger Voss", "Paul Gregu...
## $ taster_twitter_handle <chr> "@kerinokeefe", "@vossroger", "@paulgwine ...
## $ title       <chr> "Nicosia 2013 Vulkà Bianco (Etna)", "Quin...
## $ variety     <chr> "White Blend", "Portuguese Red", "Pinot Gr...
## $ winery      <chr> "Nicosia", "Quinta dos Avidagos", "Rainsto..."
```

Or a skim

```
#install.packages("skimr")
library(skimr)
skim(wine)
```




```
## Skim summary statistics
## n obs: 129971
## n variables: 14
##
```

```
## — Variable type:character —
```

variable	missing	complete	n	min	max	empty	n_unique
country	63	129908	129971	2	22	0	43
description	0	129971	129971	20	829	0	119955
designation	37465	92506	129971	1	95	0	37979
province	63	129908	129971	3	31	0	425
region_1	21247	108724	129971	3	50	0	1229
region_2	79460	50511	129971	4	17	0	17
taster_name	26244	103727	129971	10	18	0	19
taster_twitter_handle	31213	98758	129971	6	16	0	15
title	0	129971	129971	12	136	0	118840
variety	1	129970	129971	4	35	0	707
winery	0	129971	129971	1	54	0	16757

```
## — Variable type:numeric —
```

variable	missing	complete	n	mean	sd	p0	p25	p50
points	0	129971	129971	88.45	3.04	80	86	88
price	8996	120975	129971	35.36	41.02	4	17	25
x1	0	129971	129971	64985	37519.54	0	32492.5	64985

p75 p100 hist
 91 100 
 42 3300 
 97477.5 129970 

The Pipe

Starting with a dataset, you can use the pipe operator to perform operations in series.

For instance, I can filter the dataset to only show wines from Oregon

```
wine %>%
  filter(province=="Oregon")

## # A tibble: 5,373 x 14
##       X1 country description designation points price province region_1
##   <dbl> <chr>   <chr>         <chr>         <dbl> <dbl> <chr>   <chr>
## 1     2 US      Tart and s... <NA>           87    14 Oregon Willame...
## 2     4 US      Much like ... Vintner's ...  87    65 Oregon Willame...
## 3    21 US      A sleek mi... <NA>           87    20 Oregon Oregon
## 4    35 US      As with ma... Hyland         86    50 Oregon McMinnv...
## 5    41 US      A stiff, t... <NA>           86    22 Oregon Willame...
## 6    78 US      Some rosés... Rosé of        86    25 Oregon Eola-Am...
## 7   173 US      This wine ... <NA>           91    38 Oregon Willame...
## 8   233 US      There is a... Reserve        85    28 Oregon Willame...
## 9   248 US      This seems... Estate Sin...  85    45 Oregon Willame...
## 10  251 US      Spicy and ... Papillon E...  85    22 Oregon Willame...
## # ... with 5,363 more rows, and 6 more variables: region_2 <chr>,
## #   taster_name <chr>, taster_twitter_handle <chr>, title <chr>,
## #   variety <chr>, winery <chr>
```

Multiple Filters

...or I can filter the data so that I only have wines from Oregon that are over \$100

```
wine %>%
  filter(province=="Oregon") %>%
  filter(price > 100)
```



```
## # A tibble: 37 x 14
##       X1 country description designation points price province region_1
##   <dbl> <chr>   <chr>         <chr>         <dbl> <dbl> <chr>   <chr>
## 1   778 US      The Winder... Winderlea ...    92   105 Oregon Dundee ...
## 2  1082 US      Lighter in... The Tribe ...    94   120 Oregon Walla W...
## 3   5218 US      Just one b... Select Who...    92   150 Oregon McMinnv...
## 4 16333 US      From selec... Pas de Nom     94   125 Oregon Willame...
## 5 16531 US      This is a ... <NA>          96   240 Oregon Walla W...
## 6 16535 US      The aromas... Sur Echala...   95   120 Oregon Walla W...
## 7 18000 US      Focused an... Abetina        94   105 Oregon Willame...
## 8 33791 US      Focused an... Abetina        94   105 Oregon Willame...
## 9 33811 US      Multiple v... Récolte Gr...   93   125 Oregon Dundee ...
## 10 34617 US      Just two b... Olson Esta...   91   125 Oregon Dundee ...
## # ... with 27 more rows, and 6 more variables: region_2 <chr>,
## #   taster_name <chr>, taster_twitter_handle <chr>, title <chr>,
## #   variety <chr>, winery <chr>
```

Exercise

Find only the wines with variety “Pinot Gris” that cost less than \$10

Answer

```
wine %>%
  filter(variety=="Pinot Gris") %>%
  filter(price<10)

## # A tibble: 7 x 14
##       X1 country description designation points price province region_1
##   <dbl> <chr>   <chr>         <chr>      <dbl> <dbl> <chr>   <chr>
## 1  27397 US      An enjoyab... <NA>        86     9 Washing... Columbi...
## 2  28038 Moldova There's a ... Golden Lan...  86     9 Moldova  <NA>
## 3  35035 Moldova This Moldo... <NA>        86     9 Moldova  <NA>
## 4  62157 US      Firm and f... <NA>        90     9 Washing... Rattles...
## 5  64951 Argent... Flat on th... <NA>        83     9 Mendoza... Mendoza
## 6  70018 US      Light and ... <NA>        84     9 Washing... Yakima ...
## 7 117199 Romania Aromas of ... Vine in Fl...  88     9 Dealu M... <NA>
## # ... with 6 more variables: region_2 <chr>, taster_name <chr>,
## #   taster_twitter_handle <chr>, title <chr>, variety <chr>, winery <chr>
```

why do I put quotes around "Pinot Gris" but not around the price?

Combining functions

You can also combine functions using the pipe operator.

```
wine %>%
  filter(variety=="Chardonnay") %>%
  filter(province=="Oregon") %>%
  arrange(desc(points), price)

## # A tibble: 498 x 14
##       X1 country description designation points price province region_1
##       <dbl> <chr>   <chr>         <chr>         <dbl> <dbl> <chr>   <chr>
##  1 102489 US      Even if wi... Estate           96    27 Oregon Dundee ...
##  2  47900 US      This sensa... Aurora Vin...   96    60 Oregon Willame...
##  3  47902 US      This is th... Récolte Gr...   96   125 Oregon Dundee ...
##  4  31421 US      As fabulou... Shea Viney...   95    35 Oregon Willame...
##  5  78307 US      Here is an... Shea Viney...   95    35 Oregon Willame...
##  6  48377 US      Stunning i... <NA>           95    42 Oregon Willame...
##  7  47917 US      Sourced fr... Original V...   95    45 Oregon Dundee ...
##  8  95058 US      One large ... Essence         95    45 Oregon Ribbon ...
##  9  76390 US      This conve... <NA>           95    49 Oregon Willame...
## 10  31413 US      If anyone ... Hyland Vin...   95    50 Oregon McMinnv...
## # ... with 488 more rows, and 6 more variables: region_2 <chr>,
## #   taster_name <chr>, taster_twitter_handle <chr>, title <chr>,
## #   variety <chr>, winery <chr>
```

Filtered summary of select variables

Let's summarize the data from the last slide by combining filter with skim and select

```
wine %>%  
  filter(variety=="Chardonnay") %>%  
  filter(province=="Oregon") %>%  
  select(price, points) %>%  
  skim()
```



```
## Skim summary statistics
```

```
## n obs: 498
```

```
## n variables: 2
```

```
##
```

```
## — Variable type:numeric —
```

##	variable	missing	complete	n	mean	sd	p0	p25	p50	p75	p100	hist
##	points	0	498	498	89.72	2.92	80	88	90	92	96	
##	price	0	498	498	34.94	18.92	7	22	30	44.25	125	

Long exercise

1. Gather into 4 teams
2. Assign one person to “drive”
3. Create a new RMD file using the default HTML output format
4. Suggest some Oregon wines in the region of...
 - Eola-Amity Hills
 - Dundee Hills
 - Chehalem Mountains
 - Umpqua Valley

Email me the resulting html file with your marked up analysis. Don't forget to justify your recommendations!

Summary

- R is a programming language (in all its glory and ugliness)
- RMarkdown makes beautiful, reproducible, data science documents
- The Tidyverse is a philosophy (and set of packages) for data science in R

Bonus

###Beautiful tables with kable

```
wine %>%  
  filter(variety=="Chardonnay") %>%  
  filter(province=="Oregon") %>%  
  arrange(desc(points), price) %>%  
  select(points, price, province, region_1, title) %>%  
  head(5) %>%  
  knitr::kable(padding=0)
```

points	price	province	region_1	title
96	27	Oregon	Dundee Hills	The Eyrie Vineyards 2014 Estate Chardonnay (Dundee Hills)
96	60	Oregon	Willamette Valley	Ponzi 2012 Aurora Vineyard Chardonnay (Willamette Valley)
96	125	Oregon	Dundee Hills	Domaine Serene 2011 Récolte Grand Cru Chardonnay (Dundee Hills)
95	35	Oregon	Willamette Valley	Shea 2014 Shea Vineyard Chardonnay (Willamette Valley)
95	35	Oregon	Willamette Valley	Shea 2013 Shea Vineyard Chardonnay (Willamette Valley)
